

Lap Sheet - 02

Sabaragamuwa University of Sri Lanka

Faculty of Computing

Department of Software Engineering

SE6103 – Parallel And Distributed System

Name : Tharsa S.

Reg. No : 19APSE4311

Academic Period : 3rd Year 2nd Semester

Due Date : 25/11/2024

```

PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker-compose
up -d
time="2024-11-25T14:58:15+05:30" level=warning msg="D:\\3.2
lectures\\parallel\\Practicel\\Distributed-programming\\docker-compose.yml:
the attribute `version` is obsolete, it will be ignored, please remove it to
avoid potential confusion"
[+] Running 4/4
  ✓ Network distributed-programming_default    Created
0.1s
  ✓ Container namenode                        Started
1.6s
  ✓ Container datanode                       Started
1.6s
  ✓ Container historyserver                  Started
1.7s
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
datanode jps
>>
Error response from daemon: container
70d5baa87889d70c5915f21fcfe91f5282aac3c7229338803c7107a398376d62 is not
running
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker inspect
datanode
>>
[
  {
    "Id":
"70d5baa87889d70c5915f21fcfe91f5282aac3c7229338803c7107a398376d62",
    "Created": "2024-11-25T09:28:15.752295939Z",
    "Path": "/entrypoint.sh",
    "Args": [
      "/run.sh"
    ],
    "State": {
      "Status": "exited",
      "Running": false,
      "Paused": false,
      "Restarting": false,
      "OOMKilled": false,
      "Dead": false,
      "Pid": 0,
      "ExitCode": 0,
      "Error": "",
      "StartedAt": "2024-11-25T09:28:17.092829443Z",
      "FinishedAt": "2024-11-25T09:28:24.607471599Z",
      "Health": {
        "Status": "unhealthy",
        "FailingStreak": 0,

```

```

        "Log": []
    },
    "Image":
"sha256:35f899bcbe9f983825a8a3bdc135ed0e8e0eaf3b58f9b08bf257b5e86bae3b47",
    "ResolvConfPath":
"/var/lib/docker/containers/70d5baa87889d70c5915f21fcfe91f5282aac3c7229338803c7107a398376d62/resolv.conf",
    "HostnamePath":
"/var/lib/docker/containers/70d5baa87889d70c5915f21fcfe91f5282aac3c7229338803c7107a398376d62/hostname",
    "HostsPath":
"/var/lib/docker/containers/70d5baa87889d70c5915f21fcfe91f5282aac3c7229338803c7107a398376d62/hosts",
    "LogPath":
"/var/lib/docker/containers/70d5baa87889d70c5915f21fcfe91f5282aac3c7229338803c7107a398376d62/70d5baa87889d70c5915f21fcfe91f5282aac3c7229338803c7107a398376d62-json.log",
    "Name": "/datanode",
    "RestartCount": 0,
    "Driver": "overlayfs",
    "Platform": "linux",
    "MountLabel": "",
    "ProcessLabel": "",
    "AppArmorProfile": "",
    "ExecIDs": null,
    "HostConfig": {
        "Binds": null,
        "ContainerIDFile": "",
        "LogConfig": {
            "Type": "json-file",
            "Config": {}
        },
        "NetworkMode": "distributed-programming_default",
        "PortBindings": {},
        "RestartPolicy": {
            "Name": "no",
            "MaximumRetryCount": 0
        },
        "AutoRemove": false,
        "VolumeDriver": "",
        "VolumesFrom": null,
        "ConsoleSize": [
            0,
            0
        ],
        "CapAdd": null,
        "CapDrop": null,

```

```
"CgroupnsMode": "host",
"Dns": null,
"DnsOptions": null,
"DnsSearch": null,
"ExtraHosts": [],
"GroupAdd": null,
"IpcMode": "private",
"Cgroup": "",
"Links": null,
"OomScoreAdj": 0,
"PidMode": "",
"Privileged": false,
"PublishAllPorts": false,
"ReadonlyRootfs": false,
"SecurityOpt": null,
"UTSMode": "",
"UsernsMode": "",
"ShmSize": 67108864,
"Runtime": "runc",
"Isolation": "",
"CpuShares": 0,
"Memory": 0,
"NanoCpus": 0,
"CgroupParent": "",
"BlkioWeight": 0,
"BlkioWeightDevice": null,
"BlkioDeviceReadBps": null,
"BlkioDeviceWriteBps": null,
"BlkioDeviceReadIOps": null,
"BlkioDeviceWriteIOps": null,
"CpuPeriod": 0,
"CpuQuota": 0,
"CpuRealtimePeriod": 0,
"CpuRealtimeRuntime": 0,
"CpusetCpus": "",
"CpusetMems": "",
"Devices": null,
"DeviceCgroupRules": null,
"DeviceRequests": null,
"MemoryReservation": 0,
"MemorySwap": 0,
"MemorySwappiness": null,
"OOMKillDisable": false,
"PidsLimit": null,
"Ulimits": null,
"CpuCount": 0,
"CpuPercent": 0,
"IOMaximumIOps": 0,
```

```

    "IOMaximumBandwidth": 0,
    "Mounts": [
      {
        "Type": "volume",
        "Source": "distributed-programming_datanode-data",
        "Target": "/hadoop/dfs/data",
        "VolumeOptions": {}
      }
    ],
    "MaskedPaths": [
      "/proc/asound",
      "/proc/acpi",
      "/proc/kcore",
      "/proc/keys",
      "/proc/latency_stats",
      "/proc/timer_list",
      "/proc/timer_stats",
      "/proc/sched_debug",
      "/proc/scsi",
      "/sys/firmware",
      "/sys/devices/virtual/powercap"
    ],
    "ReadonlyPaths": [
      "/proc/bus",
      "/proc/fs",
      "/proc/irq",
      "/proc/sys",
      "/proc/sysrq-trigger"
    ]
  },
  "GraphDriver": {
    "Data": null,
    "Name": "overlayfs"
  },
  "Mounts": [
    {
      "Type": "volume",
      "Name": "distributed-programming_datanode-data",
      "Source": "/var/lib/docker/volumes/distributed-
programming_datanode-data/_data",
      "Destination": "/hadoop/dfs/data",
      "Driver": "local",
      "Mode": "z",
      "RW": true,
      "Propagation": ""
    }
  ],
  "Config": {

```

```

    "Hostname": "70d5baa87889",
    "Domainname": "",
    "User": "",
    "AttachStdin": false,
    "AttachStdout": true,
    "AttachStderr": true,
    "ExposedPorts": {
        "9864/tcp": {}
    },
    "Tty": false,
    "OpenStdin": false,
    "StdinOnce": false,
    "Env": [
        "CORE_CONF_fs_defaultFS=hdfs://namenode:8020",
        "PATH=/opt/hadoop-
3.2.1/bin/:/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin",
        "JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64/",
        "HADOOP_VERSION=3.2.1",
        "HADOOP_URL=https://www.apache.org/dist/hadoop/common/hadoop-
3.2.1/hadoop-3.2.1.tar.gz",
        "HADOOP_HOME=/opt/hadoop-3.2.1",
        "HADOOP_CONF_DIR=/etc/hadoop",
        "MULTIHOMED_NETWORK=1",
        "USER=root",
        "HDFS_CONF_dfs_datanode_data_dir=file:///hadoop/dfs/data"
    ],
    "Cmd": [
        "/run.sh"
    ],
    "Healthcheck": {
        "Test": [
            "CMD-SHELL",
            "curl -f http://localhost:9864/ || exit 1"
        ]
    },
    "Image": "bde2020/hadoop-datanode:latest",
    "Volumes": {
        "/hadoop/dfs/data": {}
    },
    "WorkingDir": "",
    "Entrypoint": [
        "/entrypoint.sh"
    ],
    "OnBuild": null,
    "Labels": {
        "com.docker.compose.config-hash":
"568a1045cb0015bd223e5ea3416238424b7152b2ecbc163a8461da4ef6324158",
        "com.docker.compose.container-number": "1",

```

```

        "com.docker.compose.depends_on":
"namenode:service_started:false",
        "com.docker.compose.image":
"sha256:35f899bcbe9f983825a8a3bdc135ed0e8e0eaf3b58f9b08bf257b5e86bae3b47",
        "com.docker.compose.oneoff": "False",
        "com.docker.compose.project": "distributed-programming",
        "com.docker.compose.project.config_files": "D:\\3.2
lectures\\parallel\\Practicle\\Distributed-programming\\docker-compose.yml",
        "com.docker.compose.project.working_dir": "D:\\3.2
lectures\\parallel\\Practicle\\Distributed-programming",
        "com.docker.compose.service": "datanode",
        "com.docker.compose.version": "2.29.7"
    }
},
"NetworkSettings": {
    "Bridge": "",
    "SandboxID": "",
    "SandboxKey": "",
    "Ports": {},
    "HairpinMode": false,
    "LinkLocalIPv6Address": "",
    "LinkLocalIPv6PrefixLen": 0,
    "SecondaryIPAddresses": null,
    "SecondaryIPv6Addresses": null,
    "EndpointID": "",
    "Gateway": "",
    "GlobalIPv6Address": "",
    "GlobalIPv6PrefixLen": 0,
    "IPAddress": "",
    "IPPrefixLen": 0,
    "IPv6Gateway": "",
    "MacAddress": "",
    "Networks": {
        "distributed-programming_default": {
            "IPAMConfig": null,
            "Links": null,
            "Aliases": [
                "datanode",
                "datanode"
            ],
            "MacAddress": "",
            "DriverOpts": null,
            "NetworkID":
"5914f90a19bca14f3f462747e15cfe46b108f2a275c9a5b2b28023d21f7ba931",
            "EndpointID": "",
            "Gateway": "",
            "IPAddress": "",
            "IPPrefixLen": 0,

```

```

        "IPv6Gateway": "",
        "GlobalIPv6Address": "",
        "GlobalIPv6PrefixLen": 0,
        "DNSNames": [
            "datanode",
            "70d5baa87889"
        ]
    }
}
}
}
]
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker-compose
down -v
>>
time="2024-11-25T15:00:03+05:30" level=warning msg="D:\\3.2
lectures\\parallel\\Practicel\\Distributed-programming\\docker-compose.yml:
the attribute `version` is obsolete, it will be ignored, please remove it to
avoid potential confusion"
[+] Running 6/6
  ✓ Container historyserver                               Removed
10.5s
  ✓ Container datanode                                    Removed
0.1s
  ✓ Container namenode                                    Removed
10.5s
  ✓ Volume distributed-programming_namenode-data          Removed
0.0s
  ✓ Volume distributed-programming_datanode-data          Removed
0.0s
  ✓ Network distributed-programming_default               Removed
0.3s
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker-compose
up -d
>>
time="2024-11-25T15:00:26+05:30" level=warning msg="D:\\3.2
lectures\\parallel\\Practicel\\Distributed-programming\\docker-compose.yml:
the attribute `version` is obsolete, it will be ignored, please remove it to
avoid potential confusion"
[+] Running 6/6
  ✓ Network distributed-programming_default               Created
0.1s
  ✓ Volume "distributed-programming_namenode-data"        Created
0.0s
  ✓ Volume "distributed-programming_datanode-data"        Created
0.0s
  ✓ Container namenode                                    Started
1.7s

```



```

✓ Container datanode                               Started
1.7s
✓ Container historyserver                           Started
1.8s
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode hdfs dfs -put /sample.txt /input
>>
put: `/sample.txt': No such file or directory
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode hdfs dfs -mkdir -p /input
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker
cp ./sample.txt namenode:/sample.txt
Successfully copied 2.05kB to namenode:/sample.txt
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode hdfs dfs -put /sample.txt /input
2024-11-25 09:31:24,788 INFO sasl.SaslDataTransferClient: SASL encryption
trust check: localhostTrusted = false, remoteHostTrusted = false
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode hdfs dfs -ls /input
Found 1 items
-rw-r--r--   3 root supergroup      81 2024-11-25 09:31 /input/sample.txt
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode hadoop jar /opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-
mapreducer-examples-3.2.1.jar wordcount/input/output
JAR does not exist or is not a normal file: /opt/hadoop-
3.2.1/share/hadoop/mapreduce/hadoop-mapreducer-examples-3.2.1.jar
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode ls /opt/hadoop-3.2.1/share/hadoop/mapreduce/
hadoop-mapreduce-client-app-3.2.1.jar      hadoop-mapreduce-client-hs-plugins-
3.2.1.jar      hadoop-mapreduce-client-shuffle-3.2.1.jar  lib
hadoop-mapreduce-client-common-3.2.1.jar  hadoop-mapreduce-client-jobclient-
3.2.1-tests.jar hadoop-mapreduce-client-uploader-3.2.1.jar lib-examples
hadoop-mapreduce-client-core-3.2.1.jar    hadoop-mapreduce-client-jobclient-
3.2.1.jar      hadoop-mapreduce-examples-3.2.1.jar      sources
hadoop-mapreduce-client-hs-3.2.1.jar      hadoop-mapreduce-client-nativetask-
3.2.1.jar      jdiff
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode hadoop jar /opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-
examples-3.2.1.jar wordcount /input /output
2024-11-25 09:38:50,567 INFO impl.MetricsConfig: Loaded properties from
hadoop-metrics2.properties
2024-11-25 09:38:50,621 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot
period at 10 second(s).
2024-11-25 09:38:50,621 INFO impl.MetricsSystemImpl: JobTracker metrics system
started
2024-11-25 09:38:50,869 INFO input.FileInputFormat: Total input files to
process : 1
2024-11-25 09:38:50,895 INFO mapreduce.JobSubmitter: number of splits:1

```

```
2024-11-25 09:38:51,066 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_local1344267800_0001
2024-11-25 09:38:51,067 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-11-25 09:38:51,189 INFO mapreduce.Job: The url to track the job:
http://localhost:8080/
2024-11-25 09:38:51,190 INFO mapreduce.Job: Running job:
job_local1344267800_0001
2024-11-25 09:38:51,192 INFO mapred.LocalJobRunner: OutputCommitter set in
config null
2024-11-25 09:38:51,200 INFO output.FileOutputCommitter: File Output Committer
Algorithm version is 2
2024-11-25 09:38:51,201 INFO output.FileOutputCommitter: FileOutputCommitter
skip cleanup _temporary folders under output directory:false, ignore cleanup
failures: false
2024-11-25 09:38:51,201 INFO mapred.LocalJobRunner: OutputCommitter is
org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2024-11-25 09:38:51,239 INFO mapred.LocalJobRunner: Waiting for map tasks
2024-11-25 09:38:51,240 INFO mapred.LocalJobRunner: Starting task:
attempt_local1344267800_0001_m_000000_0
2024-11-25 09:38:51,260 INFO output.FileOutputCommitter: File Output Committer
Algorithm version is 2
2024-11-25 09:38:51,261 INFO output.FileOutputCommitter: FileOutputCommitter
skip cleanup _temporary folders under output directory:false, ignore cleanup
failures: false
2024-11-25 09:38:51,281 INFO mapred.Task: Using
ResourceCalculatorProcessTree : [ ]
2024-11-25 09:38:51,285 INFO mapred.MapTask: Processing split:
hdfs://namenode:8020/input/sample.txt:0+81
2024-11-25 09:38:51,341 INFO mapred.MapTask: (EQUATOR) 0 kvi
26214396(104857584)
2024-11-25 09:38:51,341 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2024-11-25 09:38:51,341 INFO mapred.MapTask: soft limit at 83886080
2024-11-25 09:38:51,341 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2024-11-25 09:38:51,341 INFO mapred.MapTask: kvstart = 26214396; length =
6553600
2024-11-25 09:38:51,346 INFO mapred.MapTask: Map output collector class =
org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2024-11-25 09:38:51,378 INFO sasl.SaslDataTransferClient: SASL encryption
trust check: localhostTrusted = false, remoteHostTrusted = false
2024-11-25 09:38:51,484 INFO mapred.LocalJobRunner:
2024-11-25 09:38:51,486 INFO mapred.MapTask: Starting flush of map output
2024-11-25 09:38:51,486 INFO mapred.MapTask: Spilling map output
2024-11-25 09:38:51,486 INFO mapred.MapTask: bufstart = 0; bufend = 134;
bufvoid = 104857600
2024-11-25 09:38:51,486 INFO mapred.MapTask: kvstart = 26214396(104857584);
kvend = 26214348(104857392); length = 49/6553600
2024-11-25 09:38:51,496 INFO mapred.MapTask: Finished spill 0
```

```
2024-11-25 09:38:51,508 INFO mapred.Task:
Task:attempt_local1344267800_0001_m_000000_0 is done. And is in the process of
committing
2024-11-25 09:38:51,512 INFO mapred.LocalJobRunner: map
2024-11-25 09:38:51,512 INFO mapred.Task: Task
'attempt_local1344267800_0001_m_000000_0' done.
2024-11-25 09:38:51,520 INFO mapred.Task: Final Counters for
attempt_local1344267800_0001_m_000000_0: Counters: 24
    File System Counters
        FILE: Number of bytes read=316693
        FILE: Number of bytes written=842115
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=81
        HDFS: Number of bytes written=0
        HDFS: Number of read operations=5
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=1
        HDFS: Number of bytes read erasure-coded=0
    Map-Reduce Framework
        Map input records=1
        Map output records=13
        Map output bytes=134
        Map output materialized bytes=166
        Input split bytes=102
        Combine input records=13
        Combine output records=13
        Spilled Records=13
        Failed Shuffles=0
        Merged Map outputs=0
        GC time elapsed (ms)=0
        Total committed heap usage (bytes)=273154048
    File Input Format Counters
        Bytes Read=81
2024-11-25 09:38:51,521 INFO mapred.LocalJobRunner: Finishing task:
attempt_local1344267800_0001_m_000000_0
2024-11-25 09:38:51,521 INFO mapred.LocalJobRunner: map task executor complete.
2024-11-25 09:38:51,524 INFO mapred.LocalJobRunner: Waiting for reduce tasks
2024-11-25 09:38:51,525 INFO mapred.LocalJobRunner: Starting task:
attempt_local1344267800_0001_r_000000_0
2024-11-25 09:38:51,532 INFO output.FileOutputCommitter: File Output Committer
Algorithm version is 2
2024-11-25 09:38:51,532 INFO output.FileOutputCommitter: FileOutputCommitter
skip cleanup_temporary folders under output directory:false, ignore cleanup
failures: false
2024-11-25 09:38:51,533 INFO mapred.Task: Using
ResourceCalculatorProcessTree : [ ]
```

```
2024-11-25 09:38:51,538 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin:
org.apache.hadoop.mapreduce.task.reduce.Shuffle@4e1f484
2024-11-25 09:38:51,541 WARN impl.MetricsSystemImpl: JobTracker metrics system
already initialized!
2024-11-25 09:38:51,559 INFO reduce.MergeManagerImpl: MergerManager:
memoryLimit=627939712, maxSingleShuffleLimit=156984928,
mergeThreshold=414440224, ioSortFactor=10, memToMemMergeOutputsThreshold=10
2024-11-25 09:38:51,562 INFO reduce.EventFetcher:
attempt_local1344267800_0001_r_000000_0 Thread started: EventFetcher for
fetching Map Completion Events
2024-11-25 09:38:51,590 INFO reduce.LocalFetcher: localfetcher#1 about to
shuffle output of map attempt_local1344267800_0001_m_000000_0 decomp: 162 len:
166 to MEMORY
2024-11-25 09:38:51,595 INFO reduce.InMemoryMapOutput: Read 162 bytes from
map-output for attempt_local1344267800_0001_m_000000_0
2024-11-25 09:38:51,598 INFO reduce.MergeManagerImpl: closeInMemoryFile ->
map-output of size: 162, inMemoryMapOutputs.size() -> 1, commitMemory -> 0,
usedMemory ->162
2024-11-25 09:38:51,600 INFO reduce.EventFetcher: EventFetcher is interrupted.
Returning
2024-11-25 09:38:51,601 INFO mapred.LocalJobRunner: 1 / 1 copied.
2024-11-25 09:38:51,602 INFO reduce.MergeManagerImpl: finalMerge called with 1
in-memory map-outputs and 0 on-disk map-outputs
2024-11-25 09:38:51,610 INFO mapred.Merger: Merging 1 sorted segments
2024-11-25 09:38:51,610 INFO mapred.Merger: Down to the last merge-pass, with
1 segments left of total size: 148 bytes
2024-11-25 09:38:51,611 INFO reduce.MergeManagerImpl: Merged 1 segments, 162
bytes to disk to satisfy reduce memory limit
2024-11-25 09:38:51,612 INFO reduce.MergeManagerImpl: Merging 1 files, 166
bytes from disk
2024-11-25 09:38:51,614 INFO reduce.MergeManagerImpl: Merging 0 segments, 0
bytes from memory into reduce
2024-11-25 09:38:51,614 INFO mapred.Merger: Merging 1 sorted segments
2024-11-25 09:38:51,615 INFO mapred.Merger: Down to the last merge-pass, with
1 segments left of total size: 148 bytes
2024-11-25 09:38:51,616 INFO mapred.LocalJobRunner: 1 / 1 copied.
2024-11-25 09:38:51,646 INFO Configuration.deprecation: mapred.skip.on is
deprecated. Instead, use mapreduce.job.skiprecords
2024-11-25 09:38:51,672 INFO sasl.SaslDataTransferClient: SASL encryption
trust check: localHostTrusted = false, remoteHostTrusted = false
2024-11-25 09:38:51,715 INFO mapred.Task:
Task:attempt_local1344267800_0001_r_000000_0 is done. And is in the process of
committing
2024-11-25 09:38:51,718 INFO mapred.LocalJobRunner: 1 / 1 copied.
2024-11-25 09:38:51,718 INFO mapred.Task: Task
attempt_local1344267800_0001_r_000000_0 is allowed to commit now
2024-11-25 09:38:51,733 INFO output.FileOutputCommitter: Saved output of task
'attempt_local1344267800_0001_r_000000_0' to hdfs://namenode:8020/output
```

```
2024-11-25 09:38:51,734 INFO mapred.LocalJobRunner: reduce > reduce
2024-11-25 09:38:51,735 INFO mapred.Task: Task
'attempt_local1344267800_0001_r_000000_0' done.
2024-11-25 09:38:51,735 INFO mapred.Task: Final Counters for
attempt_local1344267800_0001_r_000000_0: Counters: 30
  File System Counters
    FILE: Number of bytes read=317057
    FILE: Number of bytes written=842281
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=81
    HDFS: Number of bytes written=108
    HDFS: Number of read operations=10
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=3
    HDFS: Number of bytes read erasure-coded=0
  Map-Reduce Framework
    Combine input records=0
    Combine output records=0
    Reduce input groups=13
    Reduce shuffle bytes=166
    Reduce input records=13
    Reduce output records=13
    Spilled Records=13
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=8
    Total committed heap usage (bytes)=274202624
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Output Format Counters
    Bytes Written=108
2024-11-25 09:38:51,736 INFO mapred.LocalJobRunner: Finishing task:
attempt_local1344267800_0001_r_000000_0
2024-11-25 09:38:51,736 INFO mapred.LocalJobRunner: reduce task executor
complete.
2024-11-25 09:38:52,196 INFO mapreduce.Job: Job job_local1344267800_0001
running in uber mode : false
2024-11-25 09:38:52,198 INFO mapreduce.Job: map 100% reduce 100%
2024-11-25 09:38:52,200 INFO mapreduce.Job: Job job_local1344267800_0001
completed successfully
```

2024-11-25 09:38:52,207 INFO mapreduce.Job: Counters: 36

File System Counters

FILE: Number of bytes read=633750
FILE: Number of bytes written=1684396
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=162
HDFS: Number of bytes written=108
HDFS: Number of read operations=15
HDFS: Number of large read operations=0
HDFS: Number of write operations=4
HDFS: Number of bytes read erasure-coded=0

Map-Reduce Framework

Map input records=1
Map output records=13
Map output bytes=134
Map output materialized bytes=166
Input split bytes=102
Combine input records=13
Combine output records=13
Reduce input groups=13
Reduce shuffle bytes=166
Reduce input records=13
Reduce output records=13
Spilled Records=26
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=8
Total committed heap usage (bytes)=547356672

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Input Format Counters

Bytes Read=81

File Output Format Counters

Bytes Written=108

PS D:\3.2 lectures\parallel\Practical\Distributed-programming> docker exec -it namenode hadoop jar /opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.1.jar wordcount /input/output

JAR does not exist or is not a normal file: /opt/hadoop-

3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.1.jar wordcount

```

PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode hadoop jar /opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-
examples-3.2.1.jar wordcount /input /output
JAR does not exist or is not a normal file: /opt/hadoop-
3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-examples-
3.2.1.jar wordcount /input /output
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode hadoop jar /opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-
examples-3.2.1.jar wordcount /input /output
JAR does not exist or is not a normal file: /opt/hadoop-
3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-examples-
3.2.1.jar wordcount /input /output
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode hadoop jar /opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-
examples-3.2.1.jar wordcount /input /output
2024-11-25 09:42:23,521 INFO impl.MetricsConfig: Loaded properties from
hadoop-metrics2.properties
2024-11-25 09:42:23,619 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot
period at 10 second(s).
2024-11-25 09:42:23,619 INFO impl.MetricsSystemImpl: JobTracker metrics system
started
org.apache.hadoop.mapred.FileAlreadyExistsException: Output directory
hdfs://namenode:8020/output already exists
    at
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat.checkOutputSpecs(FileO
utputFormat.java:164)
    at
org.apache.hadoop.mapreduce.JobSubmitter.checkSpecs(JobSubmitter.java:277)
    at
org.apache.hadoop.mapreduce.JobSubmitter.submitJobInternal(JobSubmitter.java:1
43)
        at org.apache.hadoop.mapreduce.Job$11.run(Job.java:1570)
        at org.apache.hadoop.mapreduce.Job$11.run(Job.java:1567)
        at java.security.AccessController.doPrivileged(Native Method)
        at javax.security.auth.Subject.doAs(Subject.java:422)
        at
org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:
1730)
            at org.apache.hadoop.mapreduce.Job.submit(Job.java:1567)
            at org.apache.hadoop.mapreduce.Job.waitForCompletion(Job.java:1588)
            at org.apache.hadoop.examples.WordCount.main(WordCount.java:87)
            at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
            at
sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
            at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.j
ava:43)
            at java.lang.reflect.Method.invoke(Method.java:498)

```

```

        at
org.apache.hadoop.util.ProgramDriver$ProgramDescription.invoke(ProgramDriver.java:71)
        at org.apache.hadoop.util.ProgramDriver.run(ProgramDriver.java:144)
        at org.apache.hadoop.examples.ExampleDriver.main(ExampleDriver.java:74)
        at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
        at
sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
        at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
        at java.lang.reflect.Method.invoke(Method.java:498)
        at org.apache.hadoop.util.RunJar.run(RunJar.java:323)
        at org.apache.hadoop.util.RunJar.main(RunJar.java:236)
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode ls /opt/hadoop-3.2.1/share/hadoop/mapreduce/
>>
hadoop-mapreduce-client-app-3.2.1.jar      hadoop-mapreduce-client-hs-plugins-
3.2.1.jar      hadoop-mapreduce-client-shuffle-3.2.1.jar      lib
hadoop-mapreduce-client-common-3.2.1.jar  hadoop-mapreduce-client-jobclient-
3.2.1-tests.jar hadoop-mapreduce-client-uploader-3.2.1.jar  lib-examples
hadoop-mapreduce-client-core-3.2.1.jar    hadoop-mapreduce-client-jobclient-
3.2.1.jar      hadoop-mapreduce-examples-3.2.1.jar      sources
hadoop-mapreduce-client-hs-3.2.1.jar      hadoop-mapreduce-client-nativetask-
3.2.1.jar      jdiff
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode hdfs dfs -rm -r /output
>>
Deleted /output
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode hadoop jar /opt/hadoop-3.2.1/share/hadoop/mapreduce/hadoop-mapreduce-
examples-3.2.1.jar wordcount /input /output
>>
2024-11-25 09:44:29,540 INFO impl.MetricsConfig: Loaded properties from
hadoop-metrics2.properties
2024-11-25 09:44:29,741 INFO impl.MetricsSystemImpl: Scheduled Metric snapshot
period at 10 second(s).
2024-11-25 09:44:29,742 INFO impl.MetricsSystemImpl: JobTracker metrics system
started
2024-11-25 09:44:30,002 INFO input.FileInputFormat: Total input files to
process : 1
2024-11-25 09:44:30,027 INFO mapreduce.JobSubmitter: number of splits:1
2024-11-25 09:44:30,173 INFO mapreduce.JobSubmitter: Submitting tokens for job:
job_local756343983_0001
2024-11-25 09:44:30,173 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-11-25 09:44:30,277 INFO mapreduce.Job: The url to track the job:
http://localhost:8080/

```



```
2024-11-25 09:44:30,279 INFO mapreduce.Job: Running job:
job_local756343983_0001
2024-11-25 09:44:30,279 INFO mapred.LocalJobRunner: OutputCommitter set in
config null
2024-11-25 09:44:30,288 INFO output.FileOutputCommitter: File Output Committer
Algorithm version is 2
2024-11-25 09:44:30,288 INFO output.FileOutputCommitter: FileOutputCommitter
skip cleanup _temporary folders under output directory:false, ignore cleanup
failures: false
2024-11-25 09:44:30,288 INFO mapred.LocalJobRunner: OutputCommitter is
org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2024-11-25 09:44:30,331 INFO mapred.LocalJobRunner: Waiting for map tasks
2024-11-25 09:44:30,331 INFO mapred.LocalJobRunner: Starting task:
attempt_local756343983_0001_m_000000_0
2024-11-25 09:44:30,365 INFO output.FileOutputCommitter: File Output Committer
Algorithm version is 2
2024-11-25 09:44:30,366 INFO output.FileOutputCommitter: FileOutputCommitter
skip cleanup _temporary folders under output directory:false, ignore cleanup
failures: false
2024-11-25 09:44:30,386 INFO mapred.Task: Using
ResourceCalculatorProcessTree : [ ]
2024-11-25 09:44:30,391 INFO mapred.MapTask: Processing split:
hdfs://namenode:8020/input/sample.txt:0+81
2024-11-25 09:44:30,448 INFO mapred.MapTask: (EQUATOR) 0 kvi
26214396(104857584)
2024-11-25 09:44:30,449 INFO mapred.MapTask: mapreduce.task.io.sort.mb: 100
2024-11-25 09:44:30,449 INFO mapred.MapTask: soft limit at 83886080
2024-11-25 09:44:30,449 INFO mapred.MapTask: bufstart = 0; bufvoid = 104857600
2024-11-25 09:44:30,449 INFO mapred.MapTask: kvstart = 26214396; length =
6553600
2024-11-25 09:44:30,455 INFO mapred.MapTask: Map output collector class =
org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2024-11-25 09:44:30,477 INFO sasl.SaslDataTransferClient: SASL encryption
trust check: localhostTrusted = false, remoteHostTrusted = false
2024-11-25 09:44:30,544 INFO mapred.LocalJobRunner:
2024-11-25 09:44:30,547 INFO mapred.MapTask: Starting flush of map output
2024-11-25 09:44:30,547 INFO mapred.MapTask: Spilling map output
2024-11-25 09:44:30,547 INFO mapred.MapTask: bufstart = 0; bufend = 134;
bufvoid = 104857600
2024-11-25 09:44:30,547 INFO mapred.MapTask: kvstart = 26214396(104857584);
kvend = 26214348(104857392); length = 49/6553600
2024-11-25 09:44:30,558 INFO mapred.MapTask: Finished spill 0
2024-11-25 09:44:30,568 INFO mapred.Task:
Task:attempt_local756343983_0001_m_000000_0 is done. And is in the process of
committing
2024-11-25 09:44:30,575 INFO mapred.LocalJobRunner: map
2024-11-25 09:44:30,575 INFO mapred.Task: Task
'attempt_local756343983_0001_m_000000_0' done.
```

2024-11-25 09:44:30,582 INFO mapred.Task: Final Counters for
attempt_local756343983_0001_m_000000_0: Counters: 24

File System Counters

FILE: Number of bytes read=316693
FILE: Number of bytes written=839567
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=81
HDFS: Number of bytes written=0
HDFS: Number of read operations=5
HDFS: Number of large read operations=0
HDFS: Number of write operations=1
HDFS: Number of bytes read erasure-coded=0

Map-Reduce Framework

Map input records=1
Map output records=13
Map output bytes=134
Map output materialized bytes=166
Input split bytes=102
Combine input records=13
Combine output records=13
Spilled Records=13
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=0
Total committed heap usage (bytes)=273154048

File Input Format Counters

Bytes Read=81

2024-11-25 09:44:30,582 INFO mapred.LocalJobRunner: Finishing task:
attempt_local756343983_0001_m_000000_0

2024-11-25 09:44:30,582 INFO mapred.LocalJobRunner: map task executor complete.

2024-11-25 09:44:30,586 INFO mapred.LocalJobRunner: Waiting for reduce tasks

2024-11-25 09:44:30,588 INFO mapred.LocalJobRunner: Starting task:
attempt_local756343983_0001_r_000000_0

2024-11-25 09:44:30,595 INFO output.FileOutputCommitter: File Output Committer
Algorithm version is 2

2024-11-25 09:44:30,595 INFO output.FileOutputCommitter: FileOutputCommitter
skip cleanup _temporary folders under output directory:false, ignore cleanup
failures: false

2024-11-25 09:44:30,597 INFO mapred.Task: Using
ResourceCalculatorProcessTree : []

2024-11-25 09:44:30,600 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin:
org.apache.hadoop.mapreduce.task.reduce.Shuffle@20567a1d

2024-11-25 09:44:30,602 WARN impl.MetricsSystemImpl: JobTracker metrics system
already initialized!

```
2024-11-25 09:44:30,620 INFO reduce.MergeManagerImpl: MergerManager:
memoryLimit=627939712, maxSingleShuffleLimit=156984928,
mergeThreshold=414440224, ioSortFactor=10, memToMemMergeOutputsThreshold=10
2024-11-25 09:44:30,624 INFO reduce.EventFetcher:
attempt_local756343983_0001_r_000000_0 Thread started: EventFetcher for
fetching Map Completion Events
2024-11-25 09:44:30,644 INFO reduce.LocalFetcher: localfetcher#1 about to
shuffle output of map attempt_local756343983_0001_m_000000_0 decomp: 162 len:
166 to MEMORY
2024-11-25 09:44:30,647 INFO reduce.InMemoryMapOutput: Read 162 bytes from
map-output for attempt_local756343983_0001_m_000000_0
2024-11-25 09:44:30,650 INFO reduce.MergeManagerImpl: closeInMemoryFile ->
map-output of size: 162, inMemoryMapOutputs.size() -> 1, commitMemory -> 0,
usedMemory ->162
2024-11-25 09:44:30,651 INFO reduce.EventFetcher: EventFetcher is interrupted.
Returning
2024-11-25 09:44:30,652 INFO mapred.LocalJobRunner: 1 / 1 copied.
2024-11-25 09:44:30,653 INFO reduce.MergeManagerImpl: finalMerge called with 1
in-memory map-outputs and 0 on-disk map-outputs
2024-11-25 09:44:30,660 INFO mapred.Merger: Merging 1 sorted segments
2024-11-25 09:44:30,660 INFO mapred.Merger: Down to the last merge-pass, with
1 segments left of total size: 148 bytes
2024-11-25 09:44:30,661 INFO reduce.MergeManagerImpl: Merged 1 segments, 162
bytes to disk to satisfy reduce memory limit
2024-11-25 09:44:30,662 INFO reduce.MergeManagerImpl: Merging 1 files, 166
bytes from disk
2024-11-25 09:44:30,662 INFO reduce.MergeManagerImpl: Merging 0 segments, 0
bytes from memory into reduce
2024-11-25 09:44:30,662 INFO mapred.Merger: Merging 1 sorted segments
2024-11-25 09:44:30,663 INFO mapred.Merger: Down to the last merge-pass, with
1 segments left of total size: 148 bytes
2024-11-25 09:44:30,664 INFO mapred.LocalJobRunner: 1 / 1 copied.
2024-11-25 09:44:30,693 INFO Configuration.deprecation: mapred.skip.on is
deprecated. Instead, use mapreduce.job.skiprecords
2024-11-25 09:44:30,724 INFO sasl.SaslDataTransferClient: SASL encryption
trust check: localhostTrusted = false, remoteHostTrusted = false
2024-11-25 09:44:30,769 INFO mapred.Task:
Task:attempt_local756343983_0001_r_000000_0 is done. And is in the process of
committing
2024-11-25 09:44:30,772 INFO mapred.LocalJobRunner: 1 / 1 copied.
2024-11-25 09:44:30,773 INFO mapred.Task: Task
attempt_local756343983_0001_r_000000_0 is allowed to commit now
2024-11-25 09:44:30,789 INFO output.FileOutputCommitter: Saved output of task
'attempt_local756343983_0001_r_000000_0' to hdfs://namenode:8020/output
2024-11-25 09:44:30,790 INFO mapred.LocalJobRunner: reduce > reduce
2024-11-25 09:44:30,790 INFO mapred.Task: Task
'attempt_local756343983_0001_r_000000_0' done.
```

2024-11-25 09:44:30,790 INFO mapred.Task: Final Counters for
attempt_local756343983_0001_r_000000_0: Counters: 30

File System Counters

FILE: Number of bytes read=317057
FILE: Number of bytes written=839733
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=81
HDFS: Number of bytes written=108
HDFS: Number of read operations=10
HDFS: Number of large read operations=0
HDFS: Number of write operations=3
HDFS: Number of bytes read erasure-coded=0

Map-Reduce Framework

Combine input records=0
Combine output records=0
Reduce input groups=13
Reduce shuffle bytes=166
Reduce input records=13
Reduce output records=13
Spilled Records=13
Shuffled Maps =1
Failed Shuffles=0
Merged Map outputs=1
GC time elapsed (ms)=10
Total committed heap usage (bytes)=274202624

Shuffle Errors

BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0

File Output Format Counters

Bytes Written=108

2024-11-25 09:44:30,791 INFO mapred.LocalJobRunner: Finishing task:
attempt_local756343983_0001_r_000000_0

2024-11-25 09:44:30,791 INFO mapred.LocalJobRunner: reduce task executor
complete.

2024-11-25 09:44:31,286 INFO mapreduce.Job: Job job_local756343983_0001
running in uber mode : false

2024-11-25 09:44:31,287 INFO mapreduce.Job: map 100% reduce 100%

2024-11-25 09:44:31,288 INFO mapreduce.Job: Job job_local756343983_0001
completed successfully

2024-11-25 09:44:31,294 INFO mapreduce.Job: Counters: 36

File System Counters

FILE: Number of bytes read=633750

```

FILE: Number of bytes written=1679300
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=162
HDFS: Number of bytes written=108
HDFS: Number of read operations=15
HDFS: Number of large read operations=0
HDFS: Number of write operations=4
HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=1
  Map output records=13
  Map output bytes=134
  Map output materialized bytes=166
  Input split bytes=102
  Combine input records=13
  Combine output records=13
  Reduce input groups=13
  Reduce shuffle bytes=166
  Reduce input records=13
  Reduce output records=13
  Spilled Records=26
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=10
  Total committed heap usage (bytes)=547356672
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=81
File Output Format Counters
  Bytes Written=108
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it
namenode hdfs -ls /output
/opt/hadoop-3.2.1/libexec/hadoop-functions.sh: line 2401: HDFS_-LS_USER: bad
substitution
/opt/hadoop-3.2.1/libexec/hadoop-functions.sh: line 2366: HDFS_-LS_USER: bad
substitution
ERROR: -ls is not COMMAND nor fully qualified CLASSNAME.
Usage: hdfs [OPTIONS] SUBCOMMAND [SUBCOMMAND OPTIONS]

```

OPTIONS is none or any of:

--buildpaths	attempt to add class files from build tree
--config dir	Hadoop config directory
--daemon (start status stop)	operate on a daemon
--debug	turn on shell script debug mode
--help	usage information
--hostnames list[,of,host,names]	hosts to use in worker mode
--hosts filename	list of hosts to use in worker mode
--loglevel level	set the log4j level for this command
--workers	turn on worker mode

SUBCOMMAND is one of:

Admin Commands:

cacheadmin	configure the HDFS cache
crypto	configure HDFS encryption zones
debug	run a Debug Admin to execute HDFS debug commands
dfsadmin	run a DFS admin client
dfsrouteradmin	manage Router-based federation
ec	run a HDFS ErasureCoding CLI
fsck	run a DFS filesystem checking utility
haadmin	run a DFS HA admin client
jmxget	get JMX exported values from NameNode or DataNode.
oev	apply the offline edits viewer to an edits file
oiv	apply the offline fsimage viewer to an fsimage
oiv_legacy	apply the offline fsimage viewer to a legacy fsimage
storagepolicies	list/get/set/satisfyStoragePolicy block storage policies

Client Commands:

classpath	prints the class path needed to get the hadoop jar and the required libraries
dfs	run a filesystem command on the file system
envvars	display computed Hadoop environment variables
fetchdt	fetch a delegation token from the NameNode
getconf	get config values from configuration
groups	get the groups which users belong to
lsSnapshottableDir	list all snapshottable dirs owned by the current user
snapshotDiff	diff two snapshots of a directory or diff the current directory contents with a snapshot
version	print the version

Daemon Commands:

balancer	run a cluster balancing utility
----------	---------------------------------

datanode	run a DFS datanode
dfsrouter	run the DFS router
diskbalancer	Distributes data evenly among disks on a given node
httpfs	run HttpFS server, the HDFS HTTP Gateway
journalnode	run the DFS journalnode
mover	run a utility to move block replicas across storage types
namenode	run the DFS namenode
nfs3	run an NFS version 3 gateway
portmap	run a portmap service
secondarynamenode	run the DFS secondary namenode
sps	run external storagepolicysatisfier
zkfc	run the ZK Failover Controller daemon

SUBCOMMAND may print help when invoked w/o parameters or with -h.

```
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it namenode hdfs dfs -ls /output
```

Found 2 items

```
-rw-r--r--  3 root supergroup          0 2024-11-25 09:44 /output/_SUCCESS
-rw-r--r--  3 root supergroup       108 2024-11-25 09:44 /output/part-r-00000
```

```
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker exec -it namenode hdfs dfs -cat /output/part-r-00000
```

```
2024-11-25 09:55:46,457 INFO sas1.SaslDataTransferClient: SASL encryption
trust check: localHostTrusted = false, remoteHostTrusted = false
```

```
Distributed      1
```

```
Let's    1
```

```
This     1
```

```
a        1
```

```
explore.    1
```

```
file       1
```

```
for        1
```

```
is         1
```

```
lecture.    1
```

```
programming 1
```

```
simple      1
```

```
the        1
```

```
txt        1
```

```
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> docker-compose
down
```

```
time="2024-11-25T15:31:07+05:30" level=warning msg="D:\3.2
```

```
lectures\parallel\Practicel\Distributed-programming\docker-compose.yml:
the attribute `version` is obsolete, it will be ignored, please remove it to
avoid potential confusion"
```

```
[+] Running 4/4
```

```
✓ Container historyserver                               Removed
```

```
10.5s
```

```
✓ Container datanode                                    Removed
```

```
10.6s
```

```
✓ Container namenode Removed
10.5s
✓ Network distributed-programming_default Removed
0.3s
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> git init
Initialized empty Git repository in D:/3.2
lectures/parallel/Practicel/Distributed-programming/.git/
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> git add .
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> git commit -m
"Distributed hadoop programming"
[master (root-commit) 3f71ffe] Distributed hadoop programming
 2 files changed, 80 insertions(+)
 create mode 100644 docker-compose.yml
 create mode 100644 sample.txt
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> git bramch -M
main
git: 'bramch' is not a git command. See 'git --help'.

The most similar command is
    branch
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> git branch -M
main
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> git remote add
origin https://github.com/Thars7/Distributed-Hadoop-Programming.git
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming> git push -u
origin main
Enumerating objects: 4, done.
Counting objects: 100% (4/4), done.
Delta compression using up to 8 threads
Compressing objects: 100% (4/4), done.
Writing objects: 100% (4/4), 774 bytes | 774.00 KiB/s, done.
Total 4 (delta 0), reused 0 (delta 0), pack-reused 0 (from 0)
To https://github.com/Thars7/Distributed-Hadoop-Programming.git
 * [new branch]      main -> main
branch 'main' set up to track 'origin/main'.
PS D:\3.2 lectures\parallel\Practicel\Distributed-programming>
```